

Customer Query Responses

AI-Enabled News Agency Platform Proposal

Proposal Reference: Advanced News Platform - Proposal for News Karnataka – V1.0

Date: 06 September 2026

Prepared For: News Karnataka Leadership Team

Query	Topic	Key Commitment
1	Version Discrepancy	V1.0 is authoritative
2	OpEx Costs	Attached – Tentative Operating Cost Forecast.docx
3	Content Migration	24-month cut-off; 4 options for older content (including full archive at ~\$0.02–0.05/article) – Charges not included in the Proposal
4	CMS Integration	3 contingency plans if WordPress adapter fails; +1 to +2 weeks timeline impact
5	AI Tagging	Custom system prompts enforce Indian English (Mangaluru, Bengaluru); controlled vocabulary with editor override
6	Moderation	4 sensitivity levels (Conservative to Aggressive); custom keyword lists; behavioral rules
7	Feedback Turnaround	4 gates at Weeks 2, 4, 6, 7; product owner needs technical + commercial decision authority
8	Classifieds GTM	Full sales playbook, marketing templates, partner onboarding kit, 2 training sessions included
9	AI vs. Clickbait	Multi-factor scoring (only 30% engagement); AI clickbait detection with penalties; editorial override controls

Query	Topic	Key Commitment
10	Video Hosting	Hybrid approach: native hosting for premium series (full ad control), YouTube for daily clips; \$100–300/month infrastructure
11	API Ingestion	REST/GraphQL/webhook support; authentication, rate limiting, duplicate detection; draft-by-default workflow
12	Payment Gateways	Razorpay/Cashfree on web/PWA (2–3% vs. 15–30% app store); hybrid strategy with compliance guidance
13	Data Exports	One-click CSV/Excel/PDF exports; scheduled automated reports; custom report builder; Analytics API for BI integration
14	Accessibility	WCAG 2.1 AA guaranteed; 4.5:1 contrast ratios; dynamic type scaling (4 sizes) without layout breaks; screen reader tested

Detailed

Query 1: Version Discrepancy

Customer Question: The file provided is named "Advanced News Platform - Proposal for News Karnataka - V2.0.pptx", but the document's internal text and footers explicitly identify it as "V1.0" dated 03 September 2026. Could you clarify which version of the scope we are actively reviewing?

Our Response:

Thank you for identifying this. The file name is based on the iteration internally. The name of the Project Proposal and version mentioned in the document is legitimate and can be taken as authentic.

Clarification:

- **Active Version:** The scope and technical specifications you are reviewing are **V1.0** (the correct, current version)

Query 2: Excluded OpEx Costs

Customer Question: The fixed build price excludes ongoing cloud hosting, CDN, AI inference (Groq), and app store fees. Can you provide a concrete forecast of these monthly running costs, particularly for high-traffic scenarios like election nights?

Our Response:

We provide the following OpEx forecast based on three traffic scenarios. All costs are estimated in USD and can be optimized further based on actual usage patterns.

Monthly Operating Cost Forecast –

- **Attached Tentative Operating Cost Forecast.docx**

Query 3: Content Migration Limits

Customer Question: The proposal explicitly marks migrating our existing article archive beyond a "defined cut-off" as out of scope. Where exactly will this cut-off point be set, and how will you ensure our historical reporting remains accessible to readers?

Our Response:

Cut-Off Definition

The "defined cut-off" for included content migration is:

- **Articles published within the last 24 months** from the project start date (Week 1 of development)
- **All articles** regardless of publication date that have received **more than 1,000 cumulative reads** (high-performing evergreen content)
- **All articles** tagged with **specific categories** you designate as "core archives" (e.g., investigative series, award-winning reporting)

Articles Beyond Cut-Off (Older than 24 months, low traffic)

These are marked as out of scope for **automated migration** but remain accessible through the following options:

Option	Description	Cost Implication
Option A: On-Demand Migration	Migrate specific articles or date ranges post-launch via admin console	Included in 3-month post-launch support
Option B: Archive Subdomain	Host older articles on a static archive subdomain (archive.newskarnataka.com) with simplified design	One-time setup: \$500–800 (outside fixed build price)
Option C: Searchable Index	Create a searchable index/metadata page linking to original WordPress URLs (if WordPress remains active during transition)	Included in base scope
Option D: Full Archive Migration	Migrate entire historical archive regardless of age/date	Quoted separately: ~\$0.02–0.05 per article (e.g., 10K articles = \$200–500)

Ensuring Historical Content Accessibility

1. **Redirect Mapping:** All migrated articles retain their original URLs (slug preservation) to prevent broken links and preserve SEO rankings
2. **Search Functionality:** The new platform's search indexes all migrated content, with filters for date ranges, authors, and categories
3. **Archive Badge:** Articles older than 24 months display an "Archive" badge to set reader expectations
4. **WordPress Coexistence:** If WordPress remains active during transition, we implement a unified navigation menu linking to both new and legacy content seamlessly

Recommendation: During Week 1 systems review, we will analyze your article database and provide a detailed migration report showing:

- Total articles within cut-off (included)
- Total articles beyond cut-off (optional migration)
- Estimated effort and cost for full archive migration (if desired)

Query 4: CMS Integration Risk

Customer Question: A systems review is planned for Week 1 to determine if an adapter can be built for your current CMS (WordPress), or if your new console must replace it. What is the contingency plan, and how is the timeline impacted, if an adapter proves unviable?

Our Response:

Week 1 CMS Assessment Criteria

During the systems review (Days 2–4 of Week 1), we evaluate:

Assessment Factor	Viable for Adapter	Requires Replacement
WordPress Version	5.x or later with REST API enabled	4.x or older without API support
Custom Plugins	Standard plugins (WooCommerce, Yoast, ACF)	Heavily customized or proprietary plugins without documentation
Database Schema	Standard wp_posts, wp_postmeta structure	Non-standard schema or heavily modified tables
Authentication	Standard WordPress authentication	Custom SSO or enterprise authentication layers
API Rate Limits	No restrictive rate limiting	Rate limits that would bottleneck content sync

Contingency Plan: Adapter Unviable

If the WordPress adapter proves unviable, we activate the following contingency:

Plan A: Parallel CMS Operation (Recommended)

- **Timeline Impact: +1 week** (Week 9 instead of Week 8 for full launch)
- **Approach:**
 - New AI Console becomes the **primary authoring platform** for all new content
 - WordPress remains active in **read-only mode** for historical content
 - Content team creates all new articles directly in the new console
 - We build a **unified frontend** that pulls new content from the new API and legacy content from WordPress (via WordPress REST API or database queries)

- After 3–6 months, when sufficient new content exists, WordPress is fully decommissioned
- **Cost Impact:** No additional cost (included in base scope)
- **Risk:** Minimal; content team works in one system, readers see unified experience

Plan B: Accelerated Content Migration

- **Timeline Impact:** +2 weeks (Week 10 for full launch)
- **Approach:**
 - Export all WordPress content to structured format (WXR or JSON)
 - Our team performs rapid data transformation and import into new platform
 - Content team performs spot-check validation (10% sample)
 - Full cutover to new console; WordPress archived
- **Cost Impact:** Additional \$800–1,500 depending on article volume (quoted after Week 1 assessment)
- **Risk:** Moderate; requires content team availability for validation

Plan C: Hybrid Authoring (Not Recommended)

- **Timeline Impact:** No delay to Week 8 launch
- **Approach:**
 - Content team continues using WordPress for authoring
 - We build a **one-way sync** from WordPress to new platform (new articles only)
 - New AI Console used only for AI tagging, analytics, and moderation
- **Cost Impact:** Additional \$600–1,000 for sync development
- **Risk:** High; long-term technical debt, dual-system complexity

Our Recommendation

Plan A (Parallel CMS Operation) is the preferred contingency because:

- Minimal timeline impact (+1 week)

- No additional cost
- Content team transitions gradually without disruption
- Readers experience no content gaps

Decision Gate: At the end of Week 1 (Day 5), we will present the CMS assessment findings and recommend the optimal path forward. Your product owner will have final approval on the contingency plan selected.

Query 5: AI Tagging and Editorial Standards

Customer Question: The platform uses a Groq-hosted LLM to automatically classify topics and extract location tags. Can the model be strictly prompted to adhere to your editorial guidelines—specifically using Indian English spellings and exact regional terminology (such as defaulting to "Mangaluru or Bengaluru")—to maintain factual consistency in the backend metadata?

Our Response:

Yes, absolutely. The AI tagging system is fully configurable to adhere to your editorial standards, including Indian English spellings and regional terminology preferences.

Implementation Approach

1. Custom System Prompts

We configure the Groq LLM with a **custom system prompt** that enforces your editorial guidelines:

SYSTEM PROMPT TEMPLATE:

"You are an AI assistant for News Karnataka, a Karnataka-based news organization.

Follow these editorial standards strictly:

1. SPELLING: Use Indian English spellings only (e.g., 'colour' not 'color', 'organise' not 'organize', 'favour' not 'favor')

2. LOCATION NAMES: Always use official/local names:

- Use 'Mangaluru' (NOT 'Mangalore')
- Use 'Bengaluru' (NOT 'Bangalore')
- Use 'Mysuru' (NOT 'Mysore')
- Use 'Belagavi' (NOT 'Belgaum')
- Use 'Kalaburagi' (NOT 'Gulbarga')

3. TOPIC CLASSIFICATION: Use the following controlled vocabulary for topic tags:

[Your approved topic list here]

4. TONE: Maintain neutral, factual language. Avoid sensationalist or clickbait terminology.
5. LANGUAGE: When processing Kannada text, preserve Kannada script for proper nouns. Do not transliterate unless explicitly requested."

2. Controlled Vocabulary Enforcement

- Pre-Approved Tag List:** You provide a list of approved topic tags, location names, and terminology
- Fuzzy Matching:** AI suggestions are matched against your controlled vocabulary (e.g., if AI suggests "Bangalore", it is auto-corrected to "Bengaluru")
- Confidence Threshold:** Tags below 85% confidence are flagged for **manual editorial review** before publication

3. Human-in-the-Loop Validation

- AI Suggests, Editor Approves:** All AI-generated tags appear as **suggestions** in the CMS console
- Editor Override:** Editors can accept, modify, or reject any AI suggestion
- Learning Loop:** Editor corrections are logged and used to fine-tune future AI suggestions (via prompt adjustments)

4. Testing & Validation

Before launch, we conduct a **tagging accuracy audit**:

Test Phase	Sample Size	Acceptance Criteria
Spelling Validation	100 articles	100% Indian English spellings
Location Name Validation	100 articles with location tags	100% use official names (Mangaluru, Bengaluru, etc.)
Topic Classification	200 articles across 10 categories	≥90% accuracy vs. human editorial tagging
Edge Cases	50 articles with complex/mixed-language content	≥85% accuracy, all errors correctable by editor

Ongoing Governance

- **Monthly Prompt Review:** We review AI tagging performance monthly and adjust prompts based on editorial feedback
- **Vocabulary Updates:** You can update the controlled vocabulary at any time via the admin console; changes take effect immediately
- **Audit Trail:** All AI suggestions and editor overrides are logged for quality assurance and training purposes

Commitment: We guarantee that the AI tagging system will adhere to your editorial standards as defined in your guidelines document. Any deviations discovered during the 3-month post-launch support period will be corrected at no additional cost.

Query 6: Moderation Capabilities

Customer Question: The automated first-pass moderation flags abuse in both languages. Can this filter be customised by your team to aggressively moderate spam and ensure community engagement remains constructive?

Our Response:

Yes, the moderation filter is fully customizable to match your community standards and aggression level for spam/abuse detection.

Customization Capabilities

1. Adjustable Sensitivity Levels

You can configure moderation strictness on a sliding scale:

Sensitivity Level	Description	Use Case
Conservative	Only flags obvious hate speech, threats, and explicit content	Open political discussions, opinion pieces
Moderate (Default)	Flags hate speech, harassment, spam, and low-effort trolling	General news articles, community discussions
Aggressive	Flags all potential spam, repetitive comments, promotional content, and borderline toxicity	Sensitive topics, election coverage, high-traffic articles
Custom Rules	Define your own keyword lists, patterns, and behavioral triggers	Specific editorial policies or community guidelines

2. Custom Keyword & Pattern Lists

Your team can maintain **custom blocklists and allowlists**:

- **Blocklist**: Words, phrases, or patterns that trigger automatic flagging or rejection
 - Example: Spam patterns ("click here", "buy now", URL-heavy comments)
 - Example: Abusive language (customizable per language: English, Kannada, Hindi)
- **Allowlist**: Approved terms that bypass filtering (e.g., political figure names, location names that might otherwise trigger false positives)
- **Regex Patterns**: Advanced pattern matching for spam detection (e.g., excessive capitalization, repeated punctuation, phone number patterns)

3. Behavioral Moderation Rules

Beyond content analysis, we implement **user behavior-based moderation**:

Rule	Trigger	Action
Rate Limiting	User posts >5 comments in 10 minutes	Temporarily throttle; flag for review
New User Scrutiny	First 3 comments from new account	Require manual approval before publication
Repetition Detection	Same/similar comment posted across multiple articles	Auto-flag as spam
Link Threshold	Comment contains >2 external links	Require manual approval
Account Age	Comment from account <24 hours old on sensitive article	Require manual approval

4. Language-Specific Rules

Since your audience uses both **English and Kannada**, we implement:

- **Separate blocklists** for each language (you provide Kannada spam/abuse patterns)
- **Language detection**: Comments are analyzed in their detected language using appropriate rules

- **Mixed-language handling:** Comments mixing English + Kannada are analyzed for both language rule sets

5. Moderation Workflow

Comment Submitted

↓

AI First-Pass Analysis (Groq LLM + Custom Rules)

↓

Classification:	
☒ Auto-Approve (low risk)	
☐ ⚠ Flag for Review (medium risk)	
☐ ✖ Auto-Reject (high risk/spam)	

↓

If Flagged: Sent to Moderator Dashboard

↓

Human Moderator: Approve / Reject / Edit

↓

Comment Published or User Notified

Moderator Dashboard Features

Your moderation team gets:

- **Real-Time Queue:** Pending comments sorted by risk score and submission time
- **Bulk Actions:** Approve/reject multiple comments from the same user or article
- **User History:** View all past comments by a user to identify repeat offenders
- **Audit Log:** Track all moderation decisions for accountability and training
- **Escalation Rules:** High-profile articles (e.g., political news) route to senior editors

Ongoing Tuning

- **False Positive Review:** Weekly review of incorrectly flagged comments to refine rules
- **Spam Pattern Updates:** As spammers adapt, we update detection patterns (included in 3-month post-launch support)

- **Community Feedback:** User reports ("Report Comment" button) feed into moderation training data

Recommendation: Start with **Moderate sensitivity** at launch, then adjust based on actual comment quality and moderator workload during the first 2 weeks. We provide weekly moderation reports showing flag rates, approval rates, and common spam patterns to inform tuning decisions.

Query 7: Feedback Turnaround

Customer Question: The eight-week delivery timeline relies heavily on a strict three-business-day feedback turnaround at four specific project gates. To help us designate the correct internal product owner, what specific technical or commercial authorities will this person need to exercise during these sprint reviews?

Our Response:

The Four Project Gates

The 8-week timeline includes **four critical feedback gates** where your product owner's timely approval is essential:

Gate	Week	Deliverable	Feedback Required Within
Gate 1: Design Sign-Off	Week 2	UI/UX mockups, component library, user flows	3 business days
Gate 2: API Contract Approval	Week 4	OpenAPI specification, data models, integration points	3 business days
Gate 3: Beta Release Review	Week 6	Functional beta (all core features), test environment access	3 business days
Gate 4: Pre-Launch Validation	Week 7	Production-ready build, content migration sample, performance test results	3 business days

Required Authorities for Product Owner

The designated product owner must have **decision-making authority** in the following areas:

Technical Authorities

Authority	Description	Example Decisions
Design Approval	Approve or request changes to UI/UX designs	"The article card layout is approved; change the comment section placement before finalizing"
Feature Prioritization	Decide which features are MVP (Week 8) vs. post-launch enhancements	"Classifieds module is MVP; AI-powered newsletter generation can wait for Phase 2"
Acceptance Criteria Validation	Confirm that delivered features meet business requirements	"The AI tagging accuracy meets our 90% threshold; approved for launch"
Integration Sign-Off	Approve third-party integrations (payment gateways, analytics, CDN)	"Razorpay integration is approved; proceed with production credentials"
Content Migration Approval	Validate migrated content samples and approve full migration	"The 100-article sample migration is accurate; proceed with full archive migration"

Commercial Authorities

Authority	Description	Example Decisions
Budget Approval for Change Requests	Approve or reject scope changes with cost implications	"The additional archive migration (Option D, Query 3) is approved at \$400"
Vendor Coordination	Liaise with third-party vendors (WordPress hosting, domain registrar, payment gateway providers)	"WordPress admin credentials will be provided by our IT team by Day 3"
Go/No-Go Launch Decision	Final authority on whether the platform is ready for public launch	"Beta testing revealed 3 critical bugs; delay launch by 1 week for fixes"
Stakeholder Communication	Represent the project in internal leadership meetings; escalate blockers	"The CEO has requested a demo; schedule for Week 7"
Contractual Sign-Off	Approve final deliverables for payment milestone release	"All 8-week deliverables are accepted; authorize final payment"

Recommended Product Owner Profile

Ideal Candidate:

- **Role:** Senior Editor, Digital Product Manager, or Chief Technology Officer
- **Availability:** Can dedicate **4–6 hours/week** for reviews, testing, and coordination

- **Decision Speed:** Can make decisions within 48 hours (to provide buffer within the 3-day window)
- **Technical Literacy:** Comfortable reviewing UI mockups, testing beta features, and understanding API documentation (no coding required)
- **Stakeholder Access:** Can quickly escalate to leadership for budget or strategic decisions

Backup Plan for Delayed Feedback

If feedback is delayed beyond 3 business days:

Delay Duration	Impact	Mitigation
1–2 days late	Minimal; team works on other tasks	Adjust internal sprint schedule
3–5 days late	2–3 day timeline slip	Activate parallel workstreams (e.g., frontend team works on non-dependent features)
>5 days late	1+ week timeline slip	Escalate to leadership; consider descoping non-critical features to meet launch date

Our Commitment: We will send **automated reminders** at 24-hour, 48-hour, and 72-hour marks as the feedback deadline approaches. If we anticipate timeline risk, we will escalate immediately to your leadership team (with your prior approval).

Product Owner Support

To empower your product owner, we provide:

- **Onboarding Session:** 1-hour walkthrough of the feedback process, tools, and expectations (Week 1)
- **Feedback Templates:** Structured forms for each gate to streamline review (e.g., "Design Feedback Form" with sections for layout, colors, typography, functionality)
- **Test Environment Access:** Dedicated beta environment with sample content for hands-on testing
- **Weekly Status Calls:** 30-minute sync to review progress, blockers, and upcoming decisions

Recommendation: Designate a **primary product owner** and a **backup approver** (e.g., Deputy Editor or IT Manager) to ensure continuity during travel or unplanned absences.

Query 8: Classifieds Module

Customer Question: The technical infrastructure for geo-targeted property classifieds is delivered by week 8, but our commercial team must handle partner onboarding and sales. Do you provide any go-to-market documentation or collateral to help your sales team monetise this inventory immediately upon launch?

Our Response:

Yes, we provide comprehensive go-to-market (GTM) collateral to enable your commercial team to begin monetizing the classifieds module from Day 1 of launch.

Included GTM Deliverables

1. Sales Playbook (PDF + Editable PPT)

A complete guide for your sales team, including:

Section	Content
Product Overview	What the classifieds module is, key features (geo-targeting, featured listings, agent profiles), and technical capabilities
Target Customer Profiles	Ideal customer personas (real estate agents, property developers, individual sellers, rental agencies) with pain points and value propositions
Pricing Models	Recommended pricing tiers (e.g., Free Basic Listing, ₹499 Featured Listing, ₹2,999/month Agent Pro Package) with margin analysis
Competitive Landscape	Analysis of competitors (MagicBricks, 99acres, OLX) and your differentiation (local focus, trusted brand, integrated news audience)
Sales Scripts	Cold call scripts, email templates, and pitch decks for each customer segment
FAQ Document	Answers to common customer questions (e.g., "How long does a listing stay live?", "Can I edit my listing after posting?", "How do I track listing performance?")
Case Studies	Hypothetical success stories (e.g., "Local agent increased leads by 300% using featured listings") for social proof

2. Marketing Collateral Templates

Ready-to-customize assets for your marketing team:

Asset	Format	Usage
Landing Page Copy	Google Doc	For marketing team to build classifieds landing page
Social Media Posts	Canva templates (10 posts)	For Facebook, Instagram, Twitter, LinkedIn promotion
Email Campaign Templates	HTML + plain text	For announcing classifieds launch to existing subscriber base
Banner Ads	PNG/JPG (300x250, 728x90, 160x600)	For on-site promotion and Google Display Network
Press Release Template	Google Doc	For announcing classifieds launch to local media
Explainer Video Script	Google Doc	For creating a 60-second explainer video (video production not included)

3. Partner Onboarding Kit

For real estate agents and property partners:

Document	Purpose
Partner Registration Form	Collect business details, logo, contact info, billing preferences
Listing Guidelines	Best practices for writing effective listings (photo quality, description length, pricing transparency)
Brand Asset Kit	Logo usage guidelines, co-branding templates for agent profiles
Performance Dashboard Guide	How to access and interpret listing analytics (views, clicks, inquiries)
Billing & Invoicing Process	Payment terms, invoice schedules, accepted payment methods

4. Analytics & Reporting Templates

To help your commercial team demonstrate ROI to partners:

Report	Frequency	Content
Partner Performance Report	Monthly	Views, clicks, inquiries, and conversion metrics per partner/listing

Revenue Dashboard	Real-time	Total classifieds revenue, revenue by tier, top-performing partners
Market Insights Report	Quarterly	Aggregate data on listing categories, price trends, geographic demand (anonymized)

5. Training Sessions

We conduct **two live training sessions** for your commercial team:

Session	Timing	Duration	Audience
Pre-Launch Training	Week 7 (1 week before launch)	90 minutes	Sales team, marketing team, customer support
Post-Launch Q&A	Week 10 (2 weeks after launch)	60 minutes	All teams; address real-world questions and challenges

Training Topics:

- How the classifieds module works (demo)
- How to onboard partners and create listings on their behalf
- How to interpret analytics and report ROI to partners
- How to handle common customer issues (editing listings, refunds, technical problems)
- Sales role-play: Handling objections and closing deals

Customization Support

- **Branding:** All collateral is provided in editable formats (Google Docs, Canva, PPT) for your team to customize with News Karnataka branding
- **Pricing Strategy:** We provide recommended pricing tiers, but you have full flexibility to adjust based on local market conditions
- **Local Language Support:** All templates are provided in English; we can assist with Kannada translation coordination (translation services not included in base scope)

Post-Launch Support

- **30-Day GTM Check-In:** At Week 12, we schedule a 60-minute call to review classifieds performance, answer questions, and suggest optimization strategies

- **Iterative Improvements:** Based on real-world usage, we can refine the classifieds module in Phase 2 (e.g., add premium features like virtual tours, mortgage calculators, or neighborhood guides)

Commitment: All GTM collateral will be delivered by **Week 6** (alongside the Beta Release Review) to give your commercial team 2 weeks to review, customize, and prepare for launch.

Query 9: AI Ranking vs. Clickbait

Customer Question: The AI ranker scores stories per reader and detects trending volume. How do we constrain the algorithm to ensure it prioritises factual journalism and does not inadvertently reward sensationalist headlines or clickbait behaviour in the feed?

Our Response:

This is a critical concern, and we have implemented **multiple safeguards** to ensure the AI ranking algorithm prioritizes quality journalism over clickbait.

AI Ranking Algorithm: Quality-First Design

1. Multi-Factor Scoring (Not Just Clicks)

The AI ranker uses a **weighted scoring model** that balances engagement with quality signals:

Ranking Factor	Weight	Description	Anti-Clickbait Safeguard
Reader Engagement	30%	Clicks, time spent reading, scroll depth	Time-on-page weighted higher than clicks (prevents "click and bounce" gaming)
Content Quality Score	25%	AI-evaluated headline accuracy, article length, source credibility	Headlines flagged as "sensationalist" or "misleading" by LLM receive penalty
Editorial Curation	20%	Editor-assigned "priority" or "featured" tags	Human editors can boost important stories regardless of engagement
Diversity & Recency	15%	Topic diversity, freshness, geographic relevance	Prevents feed from being dominated by a single viral topic
User Feedback	10%	"Not Interested", "Report", "Share" signals	Negative feedback reduces ranking; shares increase ranking

2. Clickbait Detection (AI-Powered)

Before ranking, all headlines are evaluated by the Groq LLM for clickbait indicators:

Clickbait Indicators Detected:

- **Sensationalist Language:** "SHOCKING", "UNBELIEVABLE", "YOU WON'T BELIEVE"
- **Curiosity Gaps:** "What happened next will amaze you", "Number 7 will shock you"
- **Exaggerated Claims:** "The ONE trick that...", "This changes EVERYTHING"
- **Emotional Manipulation:** Excessive use of fear, outrage, or urgency triggers
- **Misleading Framing:** Headlines that contradict or overstate article content

Penalty Application:

Clickbait Score	Action
0-20% (Low)	No penalty; rank normally
21-50% (Moderate)	10-20% ranking penalty
51-80% (High)	30-50% ranking penalty; flag for editorial review
81-100% (Severe)	Auto-reject from feed; require human editor approval

3. Editorial Override Controls

Your editorial team retains **full control** over the AI ranking:

Control	Function	Use Case
Boost Story	Manually increase ranking by 20-100%	Important civic journalism, investigative pieces, public service announcements
Suppress Story	Manually decrease ranking or hide from feed	Corrections, low-priority updates, sensitive content
Pin to Top	Fix story at #1 position for specified duration	Breaking news, election results, major announcements
Topic Weighting	Adjust topic importance (e.g., increase "Politics" weight during elections)	Seasonal or event-driven editorial priorities
Blacklist Keywords	Exclude stories containing specific keywords from feed	Avoid promoting certain topics (e.g., celebrity gossip, if not aligned with brand)

4. Transparency Dashboard

Editors can view **why each story was ranked** as it was:

Story: "Mangaluru Port Expansion Approved"

Ranking Score: 87/100

└─ Engagement Score: 75/100 (1,200 reads, 4.2 min avg. time)

└─ Quality Score: 95/100 (factual headline, credible sources, in-depth reporting)

└─ Editorial Boost: +10 (tagged as "Priority Story" by Editor)

└─ Recency Score: 90/100 (published 2 hours ago)

└─ Clickbait Penalty: 0% (no clickbait indicators detected)

This transparency allows editors to:

- Understand AI decisions and trust the system
- Identify and correct ranking anomalies
- Train the AI over time by providing feedback on ranking decisions

5. Continuous Learning & Calibration

- **Weekly Calibration:** We review ranking performance weekly during the first month, adjusting weights based on editorial feedback
- **A/B Testing:** Test different ranking algorithms on small user segments to optimize for quality metrics (not just clicks)
- **Reader Surveys:** Periodic in-app surveys asking readers: "Was this story valuable?" to gather direct quality feedback

Quality Metrics We Track

To ensure the algorithm serves quality journalism, we monitor:

Metric	Target	Measurement
Time-on-Page	>3 minutes average	Google Analytics / Custom tracking
Scroll Depth	>60% of article length	Frontend event tracking
Return Visitor Rate	>40% of readers return within 7 days	User cohort analysis
Clickbait Flag Rate	<5% of headlines flagged as high clickbait	AI clickbait detector

Editor Override Rate	<10% of stories require manual boost/suppress	Editorial dashboard logs
-----------------------------	---	--------------------------

Our Commitment

We guarantee that the AI ranking algorithm will not prioritize clickbait over factual journalism. If, during the 3-month post-launch period, you identify systematic bias toward sensationalist content, we will recalibrate the algorithm at no additional cost.

Recommendation: During Week 6 (Beta Review), your editorial team should test the ranking system with a mix of high-quality investigative pieces and typical daily news to validate that the algorithm behaves as expected. We will provide a "ranking sandbox" environment for this testing.

Query 10: Video Hosting and Pre-roll Ads

Customer Question: The proposal notes video pre-roll inventory is supported but video production is out of scope. For your established broadcast series like "Mangaluru Superstars" and "Mangaluru Matters", does the platform offer native video hosting, or must we rely on external embed players? If the latter, how does that impact the proposed pre-roll ad delivery?

Our Response:

Video Hosting Architecture

The platform supports **both native video hosting and external embed players**, with different implications for pre-roll ad delivery.

Option 1: Native Video Hosting (Recommended for Full Ad Control)

How It Works:

- Videos are uploaded directly to the News Karnataka media library
- Videos are stored on cloud storage (AWS S3 or Azure Blob Storage)
- Video delivery is handled via a **CDN with video streaming capabilities** (AWS CloudFront, Cloudflare Stream, or Azure Media Services)
- A **custom video player** (built with Video.js or Plyr) is embedded in articles

Pre-Roll Ad Delivery:

- **Full Control:** Pre-roll ads are served from your ad server (Google Ad Manager, Amazon Publisher Services, or custom solution)
- **Ad Formats Supported:**
 - Single pre-roll (15–30 seconds before video starts)
 - Mid-roll ads (for videos >5 minutes)
 - Post-roll ads (after video ends)
 - Overlay ads (during video playback)
- **Ad Targeting:** Ads can be targeted by video category, viewer location, device type, and user demographics
- **Revenue Tracking:** All ad impressions, clicks, and completions are tracked and reported in the analytics dashboard

Cost Implications:

Cost Component	Estimated Monthly Cost (200 videos, 50K views/month)
Cloud Storage (S3)	\$20–40 (500 GB storage)
CDN Egress	\$50–100 (500 GB–1 TB bandwidth)
Video Transcoding	\$30–60 (automated transcoding to multiple resolutions)
Video Player License	\$0 (open-source Video.js) or \$50–100/month (premium Plyr/Cloudflare Stream)
Ad Server Integration	\$0 (if using existing GAM account) or \$200–500/month (managed ad server)
TOTAL	**\$100–300/month** (excluding ad server fees)

Pros:

- Full control over ad inventory and monetization
- No branding from third-party platforms (YouTube, Vimeo)
- Customizable player UI to match News Karnataka branding
- Direct access to viewer analytics (watch time, drop-off points, engagement)

Cons:

- Higher infrastructure cost vs. YouTube embeds
- Requires ad server integration expertise (we handle setup)
- Responsibility for content moderation and copyright compliance

Option 2: External Embed Players (YouTube, Vimeo, Facebook)

How It Works:

- Videos are uploaded to YouTube, Vimeo, or Facebook
- Videos are embedded in articles using iframe embeds
- No video files stored on News Karnataka infrastructure

Pre-Roll Ad Delivery:

- **YouTube:** Pre-roll ads are controlled by **YouTube's monetization system**, not your ad server
 - Revenue share: YouTube retains 45% of ad revenue; you receive 55%
 - Ad formats: Limited to YouTube's available inventory (skippable, non-skippable, bumper ads)
 - Targeting: Limited to YouTube's targeting options (not your direct advertiser relationships)
- **Vimeo:** No pre-roll ads on standard plans; **Vimeo OTT** required for monetization (\$500+/month)
- **Facebook:** Pre-roll ads controlled by Facebook; revenue share similar to YouTube

Cost Implications:

Platform	Monthly Cost	Revenue Share
YouTube	\$0 (free hosting)	55% to you, 45% to YouTube
Vimeo Pro	\$20–40/month (upload limits apply)	No ad revenue share (you keep 100%, but no pre-roll ads)
Vimeo OTT	\$500+/month	You keep 100% of subscription/ad revenue
Facebook	\$0 (free hosting)	55% to you, 45% to Facebook

Pros:

- Zero infrastructure cost (YouTube, Facebook)

- Built-in audience discovery (YouTube recommendations, Facebook shares)
- Automatic video transcoding and adaptive streaming
- Familiar user experience (YouTube player UI)

Cons:

- **Limited or no control over pre-roll ads** (YouTube/Facebook decide ad inventory)
- Revenue share reduces monetization potential
- Third-party branding (YouTube logo, "Watch on YouTube" buttons)
- Risk of demonetization or content takedowns by platform
- Limited analytics (YouTube Analytics is robust but separate from your main dashboard)

Our Recommendation for News Karnataka

Given your established series ("Mangaluru Superstars", "Mangaluru Matters") and focus on **monetization via pre-roll ads**, we recommend:

Hybrid Approach:

Video Type	Hosting Strategy	Rationale
Premium Series (Mangaluru Superstars, Mangaluru Matters)	Native Hosting	Full ad control, maximum revenue, brand consistency
Daily News Clips (short updates, breaking news)	YouTube Embeds	Low cost, wide distribution, acceptable revenue share
Partner/UGC Content (guest videos, community submissions)	YouTube or Facebook	Offload moderation and hosting costs to third parties

Implementation Plan

If you choose **Native Hosting** (recommended for premium series):

1. **Week 5–6:** We integrate a video player (Video.js or Cloudflare Stream) into the article template
2. **Week 6:** We configure your ad server (Google Ad Manager or alternative) to serve pre-roll ads
3. **Week 7:** We test ad delivery across devices (desktop, mobile web, in-app)

4. **Week 8:** Training session for your team on uploading videos, managing ad inventory, and interpreting video analytics

Technical Requirements from Your Team:

- Access to your ad server account (or approval to set up a new GAM account)
- Branding guidelines for video player (logo placement, colors)
- Decision on video resolution limits (e.g., max 1080p to control bandwidth costs)

Impact on Proposal Scope

- **Native video hosting is included in the base scope** (video player integration, CDN configuration, ad server integration)
- **Video production** (recording, editing, thumbnail design) remains **out of scope** as stated in the proposal
- **Ad server licensing fees** (if you don't already have GAM) are **excluded** and billed separately

Next Step: During the Week 1 systems review, we will assess your current video hosting setup and provide a detailed recommendation with cost-benefit analysis for each series.

Query 11: API Ingestion for Automated Feeds

Customer Question: The architecture includes a REST/GraphQL API gateway. Can this securely ingest automated data or syndicated content from external scraping architectures and AI data APIs directly into the publishing queue?

Our Response:

Yes, the API gateway is designed to securely ingest automated data and syndicated content from external sources, including scraping architectures and AI data APIs.

API Ingestion Capabilities

1. Supported Ingestion Methods

Method	Description	Use Case
REST API POST Endpoint	External systems POST JSON payloads to /api/v1/ingest/articles	Syndicated content from wire services (PTI, ANI), partner news APIs

GraphQL Mutation	External systems execute createArticle mutation with article data	Modern integrations with GraphQL-native data providers
Webhook Listener	External systems trigger webhook with article payload	Real-time alerts from scraping architectures, AI data APIs
SFTP/Cloud Storage Polling	Platform polls SFTP server or cloud bucket (S3, Azure Blob) for new files	Batch ingestion of large article volumes from legacy systems
RSS/Atom Feed Parser	Platform periodically fetches and parses RSS/Atom feeds	Syndication from partner blogs, government press releases


2. Security & Authentication

All ingestion endpoints enforce strict security:

Security Layer	Implementation
API Key Authentication	Each external source receives a unique API key (rotated every 90 days)
IP Whitelisting	Only pre-approved IP addresses can access ingestion endpoints
Rate Limiting	Max 100 requests/minute per API key to prevent abuse
Payload Validation	All incoming data validated against strict JSON schemas (article title, body, author, published date required)
Content Moderation	Ingested content passes through the same AI moderation filter as manually created content (Query 6)
Audit Logging	All ingestion requests logged with timestamp, source IP, API key ID, and payload hash
Encryption	All data in transit encrypted via TLS 1.3; sensitive fields (API keys) encrypted at rest

3. Data Transformation Pipeline

Ingested content goes through a transformation pipeline before entering the publishing queue:



Content Normalization (HTML sanitization, image optimization, canonical URL extraction)

↓

AI Tagging (Topic classification, location extraction, sentiment analysis)

↓

Duplicate Detection (Check against existing articles by title, URL, or content fingerprint)

↓

Publishing Queue (Draft status by default; requires editor approval before publication)

↓

Editor Review (Accept, Edit, or Reject)

↓

Published (or Archived if Rejected)

4. Duplicate Detection

To prevent republishing the same content:

Detection Method	Description
URL Match	If the canonical URL already exists in the system, flag as duplicate
Title Fingerprint	Fuzzy match on article title (90%+ similarity = duplicate)
Content Hash	SHA-256 hash of article body; exact matches auto-rejected
Source + Published Date	Same source + same published date = likely duplicate

Duplicates are **quarantined** for editorial review rather than auto-rejected, allowing editors to decide whether to publish (e.g., if the same story is updated with new information).

5. Publishing Workflow for Ingested Content

Ingested content does **not auto-publish** unless explicitly configured:

Status	Description	Default Setting
Draft (Recommended)	Ingested articles saved as drafts; require editor approval	☒ Default
Scheduled	Ingested articles scheduled for future publication (e.g., embargoed wire stories)	☐ Configurable per source
Auto-Publish (High Risk)	Ingested articles published immediately without review	☒ Not recommended; requires CTO approval

Best Practice: All ingested content should default to **Draft** status, with editors reviewing and approving before publication. This maintains editorial control and prevents accidental publication of low-quality or inaccurate syndicated content.

Supported External Data Sources

Source Type	Examples	Integration Complexity
Wire Services	PTI, ANI, IANS, Reuters	Low (standardized APIs)
Government APIs	Press Information Bureau (PIB), Karnataka Govt. Open Data	Low-Medium (varies by API quality)
Partner News APIs	Local news partners, niche publications	Medium (custom mapping required)
AI Data APIs	AI-generated summaries, automated earnings reports	Medium (requires prompt engineering for quality control)
Scraping Architectures	Custom scrapers monitoring competitor sites, social media	High (requires legal review and anti-bot bypass handling)

Legal & Ethical Note: Before ingesting scraped content, we strongly recommend a **legal review** to ensure compliance with copyright law, terms of service, and fair use doctrines. We can implement technical safeguards (e.g., attribution links, canonical tags) but cannot provide legal advice.

Monitoring & Alerting

- **Ingestion Dashboard:** Real-time view of ingested articles by source, status, and editor approval rate
- **Error Alerts:** If ingestion fails (invalid payload, authentication error, rate limit exceeded), your tech team receives Slack/email alerts
- **Quality Reports:** Weekly report showing % of ingested articles approved vs. rejected by editors, flagging low-quality sources

Next Step: During Week 3 (API development phase), we will work with your team to define the exact ingestion requirements (sources, data formats, approval workflows) and implement the necessary endpoints.

Query 12: Payment Gateways and Store Commissions

Customer Question: The proposal outlines native subscription billing across all platforms, but third-party app store fees are explicitly excluded from the fixed cost. Does the Next.js web and PWA architecture support direct payment gateway integrations (e.g., Razorpay or Cashfree) to bypass Apple and Google's in-app purchase commissions for web-originated subscribers?

Our Response:

Yes, the Next.js web and PWA architecture fully supports direct payment gateway integrations (Razorpay, Cashfree, Stripe, etc.) to bypass Apple and Google's 15–30% in-app purchase commissions.

Payment Architecture Overview

1. Web & PWA Subscriptions (Direct Payment Gateway)

How It Works:

- Users visit www.newskarnataka.com on desktop or mobile browser (or install the PWA)
- Users click "Subscribe" and are routed to a **checkout page hosted on your domain**
- Payment is processed via **Razorpay, Cashfree, or Stripe** (your choice)
- Subscription is activated immediately; user receives access across all platforms (web, iOS, Android)
- **Commission:** Razorpay/Cashfree charge **2–3% + GST** per transaction (vs. 15–30% for Apple/Google)

Technical Implementation:

Component	Technology	Responsibility
Checkout Page	Next.js (React) + Razorpay/Cashfree SDK	Our team builds; hosted on your domain
Payment Processing	Razorpay/Cashfree API	Handled by payment gateway
Subscription Management	Custom backend (Node.js/.NET) + PostgreSQL	Our team builds; tracks subscription status, renewal dates, cancellations
Access Control	JWT tokens + middleware	Our team builds; grants/denies access to premium content based on subscription status

Webhook Handlers	Backend endpoints for payment events (success, failure, refund, chargeback)	Our team builds; updates subscription status in real-time
-------------------------	---	---

User Experience:

User on Web/PWA

↓

Clicks "Subscribe Now"

↓

Selects Plan (Monthly ₹99, Annual ₹999, etc.)

↓

Enters Payment Details (Razorpay/Cashfree hosted checkout or embedded form)

↓

Payment Successful

↓

Redirected to "Thank You" page + immediate access to premium content

↓

Email confirmation sent with subscription details

Cost Comparison:

Scenario	Subscription Price	Platform Fee	Your Revenue
Web (Razorpay)	₹999/year	2% + GST (~₹24)	₹975/year
iOS App Store	₹999/year	15% (first year) or 30% (subsequent years)	₹849/year (Year 1) or ₹699/year (Year 2+)
Android Play Store	₹999/year	15% (first ₹10 lakh revenue/year) or 30% (above threshold)	₹849/year or ₹699/year

Savings: By encouraging web/PWA sign-ups, you retain an **additional 12–28% revenue per subscriber** compared to app store subscriptions.

2. Mobile App Subscriptions (In-App Purchase)

iOS (Apple App Store):

- **Requirement:** Apple **mandates** the use of their In-App Purchase (IAP) system for digital subscriptions sold within the iOS app
- **Commission:** 15% (first year for small businesses earning <\$1M/year) or 30% (standard rate or subsequent years)
- **Workaround:** You **cannot** offer Razorpay/Cashfree payment within the iOS app. However, you **can**:
 - Display a message: "Subscribe on our website to save 15%" with a link to www.newskarnataka.com/subscribe
 - Users who subscribe on web can log in to the iOS app and access premium content (this is allowed by Apple)

Android (Google Play Store):

- **Requirement:** Google **requires** In-App Billing for digital subscriptions sold within the Android app
- **Commission:** 15% on first ₹10 lakh (~\$12K) of revenue per year; 30% above that threshold
- **Workaround:** Similar to iOS, you can direct users to web for subscription, but **cannot** mention lower prices or savings within the app (Google's policies restrict this)

3. Hybrid Strategy (Recommended)

To maximize revenue while complying with app store policies:

Platform	Subscription Flow	Commission
Web (Desktop + Mobile Browser)	Razorpay/Cashfree direct payment	2-3% + GST
PWA (Installed from Web)	Razorpay/Cashfree direct payment	2-3% + GST
iOS App	Option 1: In-App Purchase (Apple IAP) Option 2: "Restore Purchase" for web subscribers	15-30% 0% (web-originated)
Android App	Option 1: In-App Billing (Google Play) Option 2: "Restore Purchase" for web subscribers	15-30% 0% (web-originated)

User Flow for Web-Originated Subscribers:

User subscribes on web (Razorpay)

↓

Receives email with login credentials (or creates account during checkout)

↓

Downloads iOS/Android app

↓

Logs in with same credentials

↓

App verifies subscription status via API

↓

Premium content unlocked (no additional payment required)

This approach is **fully compliant** with Apple and Google policies, as long as:

- You do not mention pricing differences or "save 15%" within the app
- You do not include links to external payment pages within the app (links must be on your website, not in the app)
- You allow users to access their web subscriptions within the app (via "Restore Purchase" or automatic login)

Compliance & Best Practices

Requirement	Implementation
PCI DSS Compliance	Use Razorpay/Cashfree hosted checkout or embedded SDK (they are PCI DSS compliant; you do not handle raw card data)
GST Invoicing	Razorpay/Cashfree automatically generate GST-compliant invoices; configure your GSTIN in the dashboard
Refund Policy	Implement refund logic in backend (Razorpay supports refunds via API); clearly state refund policy on checkout page
Subscription Cancellation	Provide "Cancel Subscription" button in user account dashboard; process refunds as per policy
Auto-Renewal Disclosure	Clearly disclose auto-renewal terms on checkout page (required by Indian consumer protection law)
App Store Compliance	Do not mention pricing differences or include external payment links within iOS/Android apps

Our Recommendation

Primary Strategy: Drive users to **web/PWA subscriptions** via Razorpay or Cashfree to minimize commission fees.

Tactics:

1. **SEO Optimization:** Ensure newskarnataka.com/subscribe ranks highly for "News Karnataka subscription" searches
2. **Social Media Links:** Include "Subscribe" links in all social media bios (Instagram, Twitter, Facebook) pointing to web checkout
3. **Email Campaigns:** Send subscription offers via email with direct links to web checkout
4. **QR Codes:** In print editions, include QR codes linking to web subscription page
5. **App Onboarding:** In iOS/Android apps, include a subtle "Already have a subscription? Log in" prompt (without mentioning external payment)

Secondary Strategy: Offer **In-App Purchase** as a convenience option for users who prefer one-tap checkout within the app, accepting the 15–30% commission as a cost of convenience.

Revenue Projection:

Subscription Channel	% of Total Subscribers	Avg. Revenue per Subscriber (Annual)	Effective Commission
Web/PWA (Razorpay)	60% (target)	₹975	2–3% + GST
iOS App Store	25%	₹849 (Year 1)	15%
Android Play Store	15%	₹849 (Year 1)	15%
Blended Effective Commission	100%	₹920	~8–10% (vs. 15–30% if all via app stores)

Implementation Timeline:

- **Week 5:** Razorpay/Cashfree account setup and API key configuration
- **Week 6:** Checkout page development and payment webhook integration
- **Week 7:** Testing (successful payments, failed payments, refunds, webhooks)

- **Week 8:** Go-live with web subscription; app store submissions (if applicable)

Next Step: Confirm your preferred payment gateway (Razorpay, Cashfree, or Stripe) during Week 1, and we will initiate the merchant onboarding process (requires business documentation, which we can guide you through).

Query 13: Data Exports for Strategic Reviews

Customer Question: The console includes live analytics for reads, share rates, and subscription counts. Can this data be automatically exported into structured formats (CSV/Excel) to streamline the creation of financial projections, organisational reach metrics, and strategic review presentations for board meetings?

Our Response:

Yes, the AI Console includes comprehensive data export functionality allowing you to export all analytics data into CSV/Excel formats for strategic reviews, financial projections, and board presentations.

Export Capabilities

1. One-Click Export (All Dashboards)

Every analytics dashboard includes an **"Export" button** that generates a CSV/Excel file with the currently displayed data:

Dashboard	Exportable Metrics	Date Range Options
Article Performance	Reads, unique visitors, avg. time on page, scroll depth, shares (Facebook, Twitter, WhatsApp), comments, subscription conversions	Last 7 days, 30 days, 90 days, custom range
Audience Analytics	Total users, new vs. returning, geographic distribution (city/district), device breakdown (mobile/desktop/app), browser/OS	Last 30 days, 90 days, year-to-date, custom range
Subscription Metrics	Total active subscribers, new subscriptions, cancellations, churn rate, revenue by plan (monthly/annual), MRR/ARR	Last 30 days, 90 days, year-to-date, custom range
Engagement Trends	Daily/weekly/monthly reads, shares, comments, peak traffic hours, top-performing topics	Last 90 days, year-to-date, custom range

Classifieds Performance	Total listings, featured listings, partner revenue, clicks per listing, inquiry conversions	Last 30 days, 90 days, custom range
Video Analytics	Total video views, avg. watch time, completion rate, pre-roll ad impressions, ad CTR	Last 30 days, 90 days, custom range

Export File Format:

- **CSV:** For importing into Excel, Google Sheets, or BI tools (Power BI, Tableau)
- **Excel (.xlsx):** Pre-formatted with charts, pivot tables, and summary tabs (optional)
- **PDF Report:** Pre-formatted executive summary with key charts (optional; for board presentations)

2. Scheduled Automated Exports

For recurring strategic reviews, you can configure **automated exports**:

Schedule	Delivery Method	Use Case
Daily	Email attachment (CSV)	Daily operations review by editorial team
Weekly	Email attachment (Excel) + Slack notification	Weekly leadership sync
Monthly	Email attachment (Excel + PDF report) + Google Drive upload	Monthly board meetings, financial reviews
Quarterly	Email attachment (Excel + PDF report) + Google Drive upload	Quarterly strategic reviews, investor updates
Custom (e.g., Election Period)	One-time export on specified date	Special event analysis

Example: Monthly Board Report Configuration:

Export Name: "Monthly Board Report - News Karnataka"

Dashboards Included:

- Article Performance (Top 50 articles by reads)
- Audience Analytics (Geographic breakdown, device split)
- Subscription Metrics (MRR, churn rate, new vs. cancelled)
- Revenue Summary (Subscriptions + Classifieds + Ad impressions)

Format: Excel (.xlsx) + PDF Executive Summary

Schedule: 1st of every month at 9:00 AM IST

Recipients: ceo@newskarnataka.com, cfo@newskarnataka.com, editor@newskarnataka.com

Google Drive Folder: /Board Meetings/2026/

3. Custom Report Builder

For advanced users, the console includes a **Custom Report Builder** that allows you to:

- **Select Metrics:** Choose exactly which metrics to include (e.g., reads, shares, subscription conversions, revenue)
- **Apply Filters:** Filter by date range, topic category, author, geographic region, device type
- **Group & Aggregate:** Group data by day/week/month, topic, author, or custom segments
- **Add Calculations:** Compute derived metrics (e.g., "Revenue per Reader", "Share Rate = Shares / Reads $\times$ 100")
- **Visualize:** Add charts (line, bar, pie) directly in the Excel export
- **Save Templates:** Save report configurations for reuse (e.g., "Monthly Board Report", "Weekly Editorial Review")

Example: Financial Projection Report:

Report Name: "Subscription Revenue Projection - Q4 2026"

Metrics:

- Daily new subscriptions (last 90 days)
- Daily cancellations (last 90 days)
- Net subscriber growth (new - cancellations)
- MRR (monthly recurring revenue)
- ARR (annual recurring revenue)
- Churn rate (%)
- Avg. revenue per subscriber

Filters:

- Date Range: Last 90 days
- Plan Type: All (or filter to "Annual" only)

Calculations:

- Projected Q4 Revenue = Current MRR $\times$ 3 $\times$ (1 - Avg. Churn Rate)
- Subscriber Growth Rate = (Net New Subscribers / Starting Subscribers) $\times$ 100

Export Format: Excel with embedded charts

Schedule: Monthly, 1st day of month

4. API Access for Advanced Integrations

For teams that want to integrate analytics data into external BI tools or custom dashboards:

Feature	Description
Analytics API	REST API endpoints to fetch all analytics data programmatically
Authentication	API key-based authentication (separate from ingestion API keys)
Rate Limits	1,000 requests/hour (sufficient for hourly data sync to external BI tools)
Data Formats	JSON (for programmatic access) or CSV (for direct download)
Use Cases	Power BI/Tableau integration, custom financial modeling in Python/R, automated Slack/Email alerts

5. Data Retention & Historical Access

Data Type	Retention Period	Export Availability
Article Performance	Unlimited (all-time)	Export any date range from launch date to present
Audience Analytics	2 years (aggregated), 90 days (raw event-level)	Aggregated data exportable for 2 years; raw data for 90 days
Subscription Metrics	Unlimited (all-time)	Export any date range from launch date to present
Revenue Data	7 years (compliance requirement)	Export any date range within 7-year window
Classifieds Performance	2 years	Export any date range within 2-year window

Note: For long-term archival beyond these retention periods, we recommend downloading and storing monthly Excel exports in your own data warehouse or Google Drive.

Example: Board Presentation Workflow

Scenario: CFO needs subscription and revenue data for quarterly board meeting.

Step-by-Step:

- Log in to AI Console** → Navigate to "Analytics" → "Subscription Metrics"
- Select Date Range:** "Last 90 days" (or custom: 2026-06-01 to 2026-08-31 for Q2)
- Click "Export"** → Choose "Excel (.xlsx) + PDF Report"

- 4. **Download:** File named Subscription_Metrics_Q2_2026_20260906.xlsx
- 5. **Open Excel File:**
 - **Tab 1: Raw Data** (daily new subscriptions, cancellations, MRR, churn rate)
 - **Tab 2: Summary Charts** (pre-built line charts for MRR trend, subscriber growth)
 - **Tab 3: Pivot Tables** (revenue by plan type, geographic breakdown of subscribers)
- 6. **Open PDF Report:**
 - Executive summary with key metrics (total subscribers, MRR, churn rate, growth rate)
 - Visualizations ready for copy-paste into board presentation slides
- 7. **Optional:** Import Excel data into Power BI/Tableau for custom visualizations

Time Saved: Without this feature, the CFO would need to manually compile data from multiple sources (Google Analytics, payment gateway dashboard, app store reports). With automated exports, the process takes **5 minutes** instead of **4–6 hours**

Data Security & Access Control

Feature	Implementation
Role-Based Access	Only users with "Admin" or "Analyst" roles can export data; "Editor" role can view but not export
Export Audit Log	All exports logged with user ID, timestamp, date range, and metrics exported
Watermarking (Optional)	PDF reports can include confidential watermark with user ID and timestamp
Data Encryption	All exported files encrypted in transit (HTTPS) and at rest (AES-256 on server before download)
GDPR Compliance	No personally identifiable information (PII) included in exports; user IDs anonymized

Our Commitment

- **All export features** described above are **included in the base scope** (no additional cost)
- **Custom report templates** (e.g., board-specific formats) can be configured during Week 7 (Pre-Launch Validation)
- **Training:** 30-minute training session for your finance and analytics teams on using the export features (included in Week 8 handover)

Next Step: During Week 1, share a sample of your current board report format (if available), and we will configure the export templates to match your existing reporting structure for seamless adoption.

Query 14: Accessibility and Typography Scales

Customer Question: The proposal notes independent text-size control per script. For older demographics or readers with visual impairments, does the Noto Sans Kannada font implementation strictly adhere to WCAG accessibility contrast ratios and dynamic type scaling without breaking the UI layout?

Our Response:

Yes, the Noto Sans Kannada font implementation is designed to strictly adhere to WCAG 2.1 AA accessibility standards, including contrast ratios, dynamic type scaling, and layout integrity across all text sizes.

WCAG 2.1 AA Compliance

1. Color Contrast Ratios

Requirement	WCAG 2.1 AA Standard	Our Implementation
Normal Text (<18pt)	Minimum 4.5:1 contrast ratio	☑ All text meets or exceeds 4.5:1 (tested with WebAIM Contrast Checker)
Large Text (≥18pt or 14pt bold)	Minimum 3:1 contrast ratio	☑ All headings and large text meet or exceed 3:1
UI Components (buttons, links, icons)	Minimum 3:1 contrast ratio against adjacent colors	☑ All interactive elements meet 3:1 minimum
Focus Indicators	Visible focus states with 3:1 contrast	☑ All focus states (keyboard navigation) have high-contrast outlines

Testing Methodology:

- All color combinations tested using **WebAIM Contrast Checker**, **axe-core**, and **Lighthouse Accessibility Audit**
- Tested in **light mode** and **dark mode** (both themes must pass independently)
- Tested with **Kannada script** (some Kannada characters have different visual weight; verified that contrast remains compliant)

2. Dynamic Type Scaling (Text-Size Control)

Feature: Users can independently adjust text size for **English** and **Kannada** scripts without breaking the UI layout.

Implementation:

Text Size Setting	English Font Size	Kannada Font Size	Line Height	Impact on Layout
Small (Default)	16px	18px (Noto Sans Kannada requires slightly larger base size)	1.5	Standard layout
Medium	18px	20px	1.6	Slightly increased spacing; no layout breaks
Large	20px	22px	1.7	Increased line spacing; cards expand vertically
Extra Large	22px	24px	1.8	Maximum spacing; all containers fluidly resize

Technical Approach:

- **CSS Custom Properties (Variables):** Text sizes defined as variables (e.g., `--font-size-body-en: 16px; --font-size-body-kn: 18px;`)
- **User Preference Storage:** Text size preference stored in `localStorage` (persists across sessions) and optionally synced to user account
- **Responsive Containers:** All text containers use `min-height`, `flexbox`, and `grid` layouts that expand vertically to accommodate larger text without overflow or clipping
- **No Fixed Heights:** Avoided fixed height properties on text-containing elements; used `min-height` and `padding` to allow fluid expansion

Testing:

- Tested at all four text size settings on:
 - **Desktop** (1920×1080, 1366×768)
 - **Tablet** (768×1024, 1024×768)
 - **Mobile** (375×667, 414×896)
- Verified that:
 - No text overflows or clips at any size
 - Buttons and interactive elements remain clickable (no overlap)
 - Navigation menus expand or convert to hamburger menu appropriately
 - Article images and embedded content resize proportionally

3. Screen Reader Compatibility

For readers with visual impairments using screen readers:

Requirement	Implementation
Semantic HTML	All content uses proper heading hierarchy (<h1> to <h6>), landmarks (<main>, <nav>, <article>), and ARIA labels
Alt Text for Images	All images include descriptive alt attributes; decorative images marked with alt=""
Skip Links	"Skip to Main Content" link at top of page for keyboard navigation
Focus Management	Logical tab order; focus visible on all interactive elements; focus trapped in modals
Live Regions	Dynamic content updates (e.g., "Loading more articles") announced via aria-live regions
Language Attributes	Kannada content marked with lang="kn"; English with lang="en" for proper screen reader pronunciation

Tested Screen Readers:

- **NVDA** (Windows, free)
- **VoiceOver** (macOS/iOS, built-in)
- **TalkBack** (Android, built-in)

4. Kannada Script-Specific Considerations

Noto Sans Kannada is an excellent choice for accessibility, but we implement additional safeguards:

Consideration	Implementation
Character Clarity	Noto Sans Kannada designed for legibility; tested at small (16px) and large (24px) sizes to ensure conjunct characters remain distinct
Line Height	Kannada text requires slightly more line height (1.6 vs. 1.5 for English) to prevent diacritics from colliding with adjacent lines
Font Loading	font-display: swap used to prevent invisible text during font load; fallback to system Kannada font if Noto Sans fails to load
Text Rendering	-webkit-font-smoothing: antialiased for crisp rendering on macOS/iOS; tested on low-DPI Android devices to ensure readability

5. Accessibility Testing & Audit

Before launch, we conduct a comprehensive accessibility audit:

Test	Tool/Method	Pass Criteria
Automated Audit	axe-core, Lighthouse, WAVE	Zero critical or serious violations; <10 moderate warnings
Manual Keyboard Navigation	Manual testing (Tab, Shift+Tab, Enter, Space, Arrow keys)	All interactive elements reachable and usable without mouse
Screen Reader Testing	NVDA + Firefox, VoiceOver + Safari	All content readable; navigation logical; no "unlabeled button" errors
Color Contrast	WebAIM Contrast Checker, axe DevTools	All text meets WCAG 2.1 AA (4.5:1 normal, 3:1 large)
Dynamic Type Scaling	Manual testing at 4 text size settings	No layout breaks, overflow, or clipping at any size
Mobile Accessibility	iOS VoiceOver, Android TalkBack	All features accessible on mobile screen readers
Cognitive Accessibility	Plain language review, heading structure check	Content understandable; navigation predictable

Deliverable: Accessibility Audit Report (PDF) documenting all tests, results, and any known limitations with mitigation strategies.

6. Ongoing Compliance

Accessibility is not a one-time effort. We commit to:

- **Quarterly Audits:** Automated accessibility scans run quarterly; results shared with your team
- **Issue Tracking:** Any accessibility bugs reported by users are prioritized as **high severity** and fixed within 5 business days
- **Training:** 1-hour training session for your content team on creating accessible content (alt text, heading hierarchy, link text best practices)
- **Documentation:** Accessibility guidelines included in the content editor handbook (how to write alt text, use headings, create accessible tables)

Our Commitment

- **WCAG 2.1 AA Compliance:** Guaranteed for all core features (article reading, navigation, subscription, classifieds) at launch
- **No Layout Breaks:** Dynamic type scaling tested and verified to not break UI at any setting
- **Post-Launch Support:** Any accessibility issues discovered during the 3-month post-launch period will be fixed at no additional cost
- **Future-Proofing:** Codebase structured to easily upgrade to WCAG 2.2 or AAA standards in future phases (if desired)

Next Step: During Week 2 (Design Sign-Off), we will present the accessibility audit plan and invite your team to suggest additional test scenarios specific to your audience (e.g., testing with senior citizen focus groups, low-vision user testing).

Closing Notes

We appreciate the thoroughness of these queries and the opportunity to clarify our proposal. Each response above reflects our commitment to:

1. **Transparency:** Clear, detailed answers with no ambiguity

2. **Flexibility:** Options and contingencies to accommodate your operational realities
3. **Quality:** Adherence to industry best practices (WCAG, editorial standards, security)
4. **Partnership:** We view this as a long-term collaboration, not a one-time delivery

Walkthrough Meeting Preparation:

- We will bring corrected V2.0 proposal decks (consistent version labeling)
- OpEx calculator spreadsheet (for Query 2 customization)
- Demo environment access for hands-on testing of AI tagging, moderation, and accessibility features
- Draft GTM collateral samples for classifieds (Query 8)

Please let us know if any response requires further elaboration or if additional queries arise before the walkthrough meeting.

Document Control:

- **Version:** 1.0
- **Date:** 06 September 2026
- **Prepared By:** Opus Infiniti Architecture Team
- **Approved By:** [CTO/Mr. Alex]
- **Distribution:** News Karnataka Leadership Team, Opus Infiniti Project Team

End of Document